
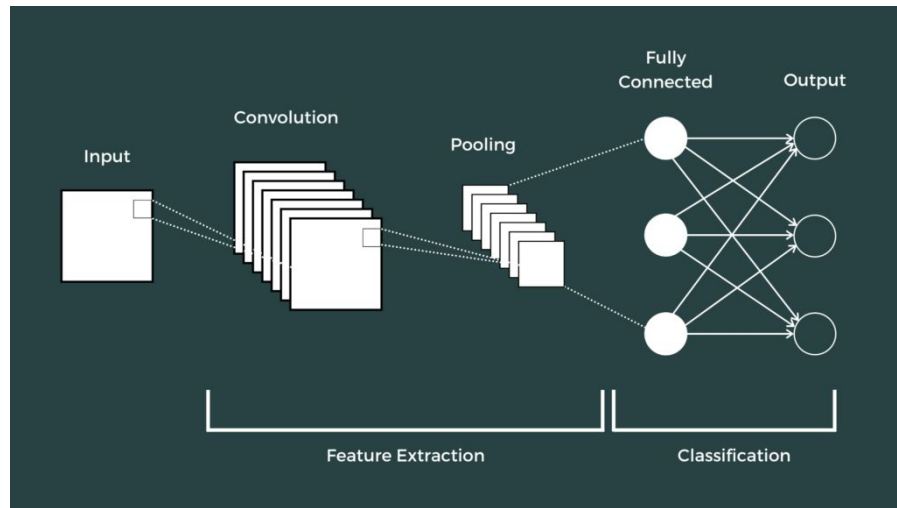


Handwritten Digits Recognition

Presented by Group 10


Mark Garcia, Alex Hwang, Hanson Nguyen

Andrew Phan, Anthony Reyes, and Linda Trinh



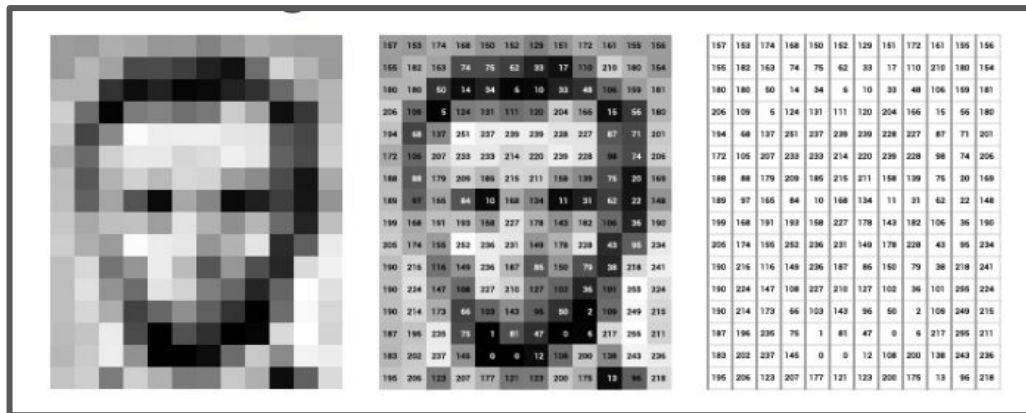


Legend:

Layer	ComponentwiseCopy	Concatenate



Data interpretation

- Machine doesn't “see” images
- Images are a large array of numbers



Data interpretation

- Conversion required for improve time and space
- Formula from Matplotlib library for conversion

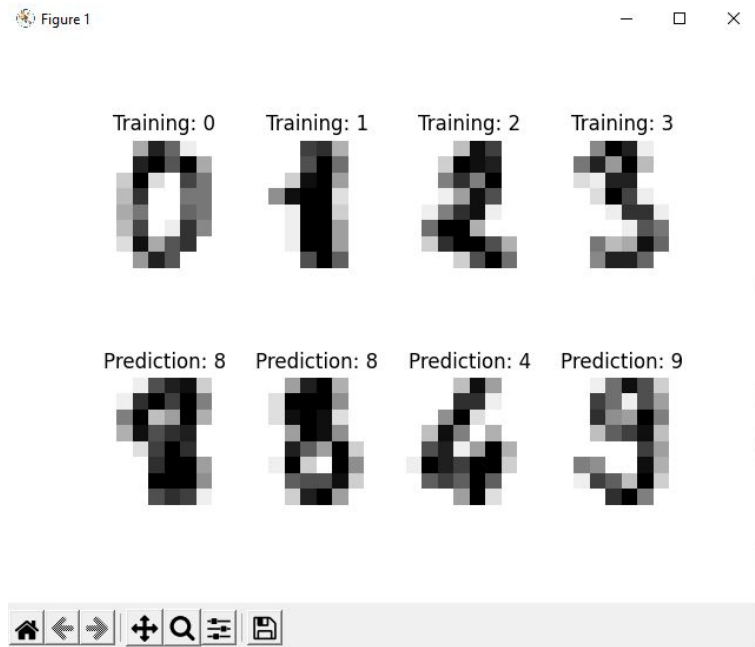
Gray Pixels = 0.2989 * R + 0.5870 * G + 0.1140 * B

- Picture converted into tensor of integers

[illegible]

Data-to-training

- Support Vector Machines classifier, or the SVM for short.
- Takes in test data and training data.
- Scans the training data numbers first, then deciphers the test data numbers.




Data-to-training

- Program creates a pixelated graph for each number.
- Pixelated graphs become training material.

Run: main


Classification report for classifier SVC(gamma=0.001):

	precision	recall	f1-score	support
0	1.00	0.99	0.99	88
1	0.99	0.97	0.98	91
2	0.99	0.99	0.99	86
3	0.98	0.87	0.92	91
4	0.99	0.96	0.97	92
5	0.95	0.97	0.96	91
6	0.99	0.99	0.99	91
7	0.96	0.99	0.97	89
8	0.94	1.00	0.97	88
9	0.93	0.98	0.95	92
accuracy			0.97	899
macro avg	0.97	0.97	0.97	899
weighted avg	0.97	0.97	0.97	899




Long Short-Term Memory (LSTM)

- Subcategory of recurrent neural network (RNN)
- Cells are most commonly composed of 3 gates (input, output, forget)
- Input data passes into one cell, is modified in the cell, and then is outputted to the next layer and also the next cell in the same layer.



LSTM (cont)

- Each cell of an LSTM layer does 4 things.
 - Decide what needs to be forgotten
 - Take in new input
 - Modify current cell state
 - Output the modified cell state



- The forget-gate decides what needs to be forgotten
- The input-gate takes in new information
- The output-gate filters the modified cell state and decides what information should be passed to the next nodes / layers

Our LSTM Model


- Comprised of 2 LSTM layers, 2 Dense Layers, and Dropout layers in between each to help with over fitting.
- Results in an accuracy rating of 98.88% with a loss of .0411

```
# RNN USING LSTM LAYERS
import numpy as np
import tensorflow as tf
from tensorflow import keras
from keras.layers import Dense, LSTM, Dropout
import cv2 as cv
import matplotlib.pyplot as plt

# import the mnist dataset that contains images of handwritten digits
(xTrain, yTrain), (xTest, yTest) = keras.datasets.mnist.load_data()

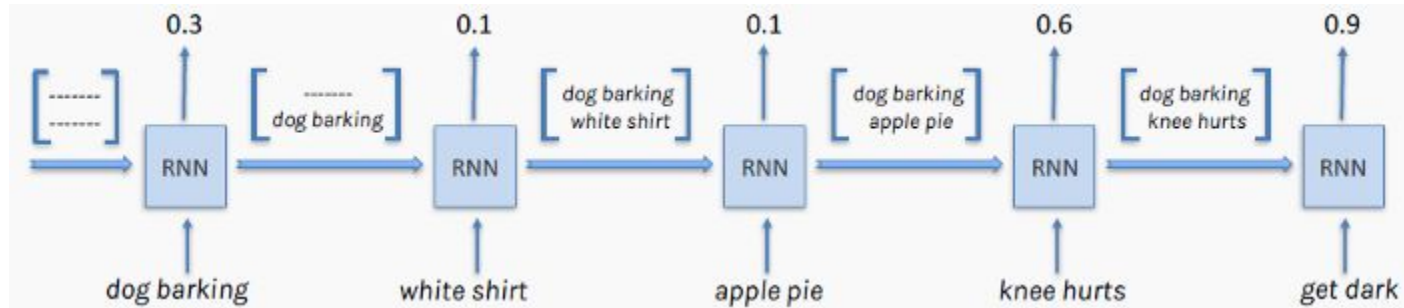
# normalize training and testing data to be btwn 0 - 1
xTrain = tf.keras.utils.normalize(xTrain, axis=1)
xTest = tf.keras.utils.normalize(xTest, axis=1)

model = tf.keras.models.Sequential()
# add LSTM layers to model followed by a dropout layer as my "output" layer
model.add(LSTM(128, input_shape=(xTrain.shape[1:]), activation='relu', return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(128, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))
# compile model
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True), optimizer='adam', metrics=['accuracy'])
# train out model with the training sets
history = model.fit(xTrain, yTrain, epochs=10, validation_data=(xTest, yTest))
```




Recurrent Neural Network (RNN)

- Artificial neural network
- Why is it important?
 - Many different use cases
 - Used in popular apps



Our RNN Model

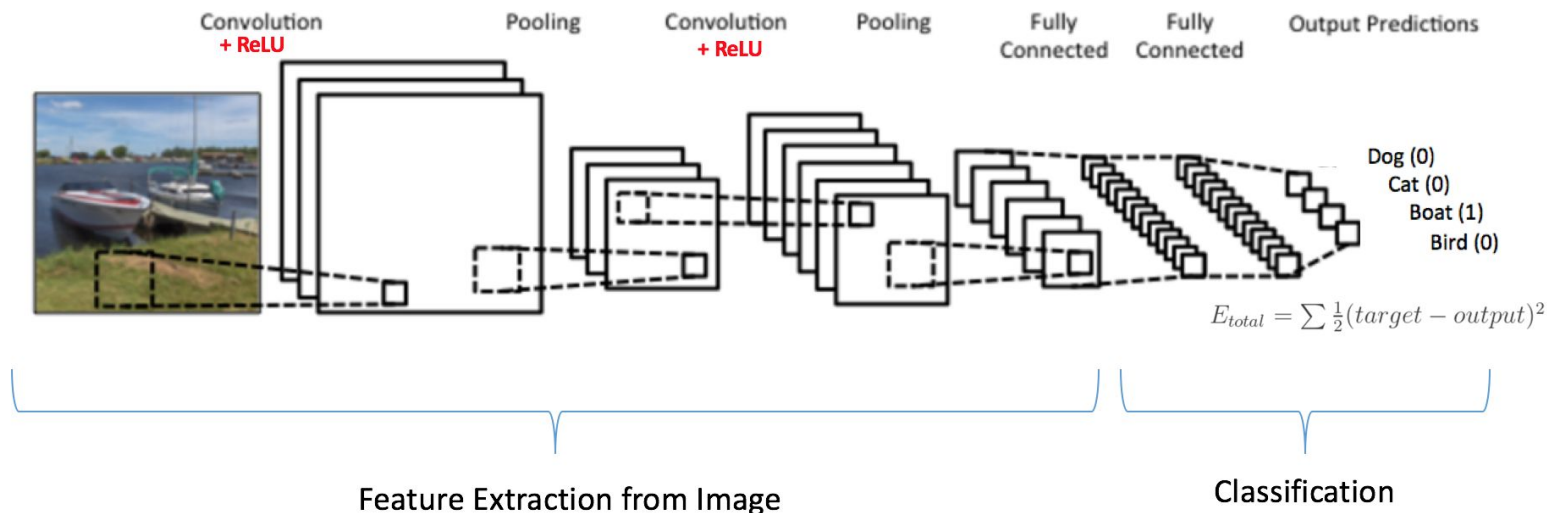
- Tensorflow
- 3 Dense layers, 1 Flatten layer
- 96.69% accuracy
- 0.1267 loss



```
database = tf.keras.datasets.mnist
(xTrain, yTrain), (xTest, yTest) = database.load_data()
xTrain = tf.keras.utils.normalize(xTrain, axis=1)
xTest = tf.keras.utils.normalize(xTest, axis=1)
#build our rnn model
model = tf.keras.models.Sequential()
# # add our layers
model.add(tf.keras.layers.Flatten(input_shape=(28, 28)))
model.add(tf.keras.layers.Dense(units=128, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(units=128, activation=tf.nn.relu))
#output layer
model.add(tf.keras.layers.Dense(units=10, activation=tf.nn.softmax))
#compile model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
epochs = 10
#train the rnn model and keep history of accuracy and loss
history = model.fit(xTrain, yTrain, epochs=epochs, validation_split=.3)
loss, accuracy = model.evaluate(xTest, yTest)
print("accuracy: ", accuracy)
print("loss: ", loss)
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs_range = range(epochs)
#plot the training and validation loss/accuracy~
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```

Convolutional Neural Network (CNN)

- What is it commonly used for?
 - Image recognition
 - Classification
- Consist of three main layers:
 - Convolutional
 - Pooling
 - Fully-connected



Our CNN Model

- Similar to RNN
 - TensorFlow
 - Split up datasets

```
# Train with convolutional neural networks (CNN)

# Import libraries
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

# Load MNIST dataset
mnist = tf.keras.datasets.mnist
# Divide training and testing datasets: training = 60,000 & testing = 10,000
(xTrain, yTrain), (xTest, yTest) = mnist.load_data()

# Pre-process & Normalize data before CNN
xTrain = tf.keras.utils.normalize(xTrain, axis = 1)
xTest = tf.keras.utils.normalize(xTest, axis = 1)
plt.imshow(xTrain[0], cmap=plt.cm.binary)

x_trainr = np.array(xTrain).reshape(-1, 28, 28, 1)
x_testr = np.array(xTest).reshape(-1, 28, 28, 1)

# Build our CNN model
model = tf.keras.models.Sequential()

# Convolution Layer 1
model.add(tf.keras.layers.Conv2D(64, (3, 3), input_shape=x_trainr.shape[1:]))
model.add(tf.keras.layers.Activation("relu"))
model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))

# Convolution Layer 2
model.add(tf.keras.layers.Conv2D(64, (3, 3)))
model.add(tf.keras.layers.Activation("relu"))
model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))

# Convolution Layer 3
model.add(tf.keras.layers.Conv2D(64, (3, 3)))
model.add(tf.keras.layers.Activation("relu"))
model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2)))

# Fully Connected Layer 1
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(64))
model.add(tf.keras.layers.Activation("relu"))
```

```

# Fully Connected Layer 2
model.add(tf.keras.layers.Dense(32))
model.add(tf.keras.layers.Activation("relu"))

# Fully Connected Layer 3
model.add(tf.keras.layers.Dense(10))
model.add(tf.keras.layers.Activation("softmax"))

# Compile the model
model.compile(loss="sparse_categorical_crossentropy", optimizer='adam', metrics=['accuracy'])
epochs = 10
# Train the model
history = model.fit(x_trainr, yTrain, epochs=epochs, validation_split=0.3)

# Evaluate the model
loss, accuracy = model.evaluate(x_testr, yTest)
print("loss: ", loss)
print("accuracy: ", accuracy)

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']


loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

```




```


loss: 0.06782113015651703
accuracy: 0.9829999804496765

```


Comparing Results



LSTM




RNN



CNN

Comparing Results

- LSTM
 - Accuracy: 98.88%
 - Loss: .0411
- CNN
 - Accuracy: 98.22%
 - Loss: .0641
- RNN
 - Accuracy: 96.69%
 - Loss: .1267



Thank you for listening!